

Neural Networks



Justin Chen, Charles Ma
Sept. 26, 2017

MIC Conference

Data: TBA


Project ideas:

Signup:

Brief Recap:


- Data-driven learning
-

What's a Neural Network?

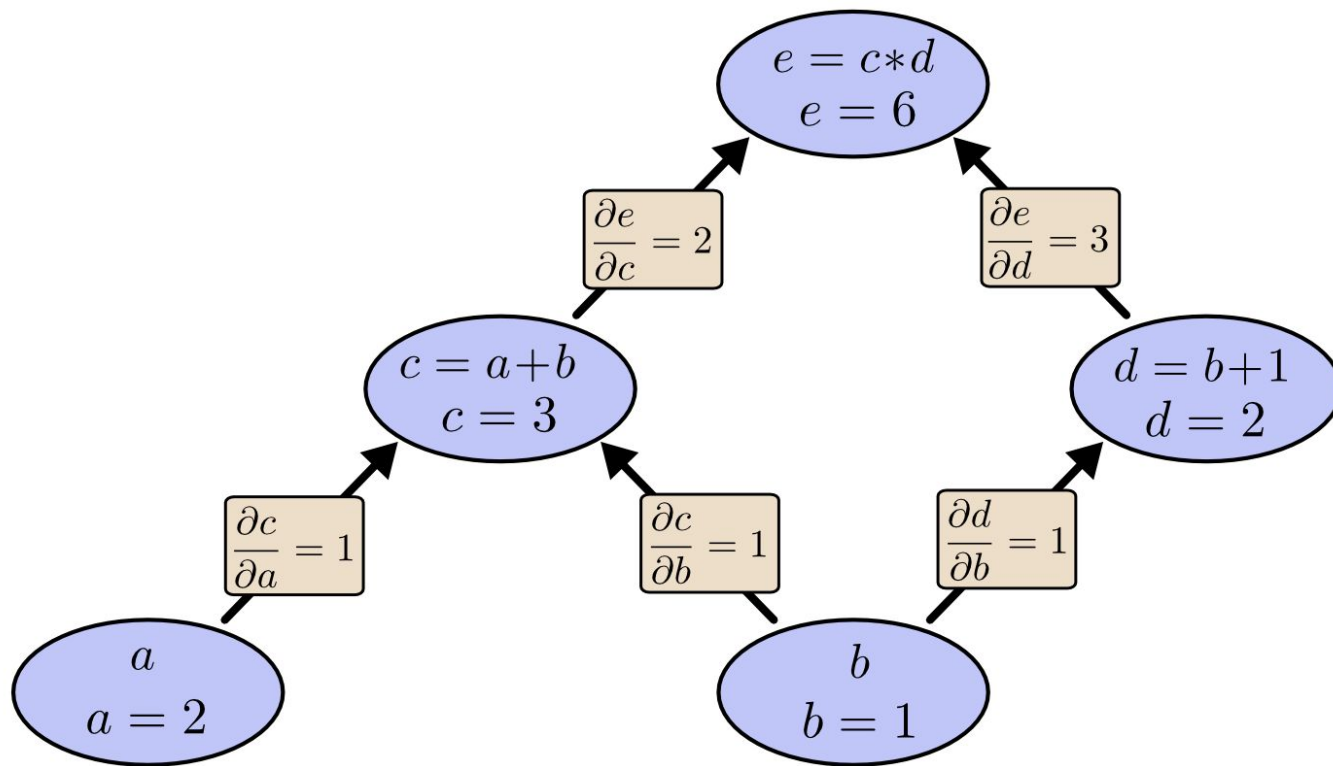


Vectors and Vectorization


- Data is represented in arrays, with each term representing one dimension
- Vectorization is essentially using built-in functions to take advantage of parallel computing capabilities (reduces training time)





Computation Graph



Neural Network (who decides function f)




Neural Networks as Computational Graphs



Gradient Descent

- The process by which your neural net is trained to be accurate
- Through numerous
Iterations of logistic
regression
- Reach the global minima!
-




Logistic Regression Cost Function

- Used when output of neural network is 0 (false) or 1 (true)

Activation Functions


ReLU

(Rectified Linear Unit):




$$\max(0, x)$$

Sigmoid (σ):



$$1/(1+e^{-x})$$

Tanh:



$$2*\sigma(2x)-1$$

Levels of Supervision During Learning

Supervised Learning - you *have* labeled data; your data comes in pairs of input and desired output.

Unsupervised Learning - you *do not have* labeled data; your neural net looks to group different inputs based on similarities i.e. *clustering*.

Semi-supervised Learning - you have *some* labeled data; labeled data is expensive to generate, so you use a mix of labeled and unlabelled.

A Tidbit on the Universal Approximation Theorem

A multilayer perceptron can approximate *continuous functions* on a compact subset of real numbers with mild assumptions on the activation function. (Thanks Wikipedia)

A Basic Multilayer Perceptron with PyTorch

```
import torch
import torch.nn as nn

class MLP(nn.Module):

    # one hidden layer
    def __init__(self, num_inputs, num_hiddens):
        super(MLP, self).__init__()
        self.fc1 = nn.Linear(num_inputs, num_hiddens)
        self.relu = nn.ReLU()
        self.fc2 = nn.Linear(num_hiddens, 1)

    def forward(self, x):
        output = self.fc1(x)
        output = self.relu(output)
        output = self.fc2(output)
        return output

net = MLP(5, 3)
print(net)

parameters = list(net.parameters())
print(parameters)
print(len(parameters))
print(parameters[0].size())
```

Output:

```
MLP (
  (fc1): Linear (5 -> 3)
  (relu): ReLU ()
  (fc2): Linear (3 -> 1)
)
[Parameter containing:
-0.1721  0.0375 -0.4098  0.0769 -0.4214
 0.0392 -0.1720  0.3798  0.1420 -0.4239
 0.3698  0.3336 -0.4313  0.3198  0.2039]
[torch.FloatTensor of size 3x5]
, Parameter containing:
-0.2104
 0.0990
 0.1316]
[torch.FloatTensor of size 3]
, Parameter containing:
-0.2340 -0.2816  0.4150]
[torch.FloatTensor of size 1x3]
, Parameter containing:
 0.1908]
[torch.FloatTensor of size 1]
]
4
torch.Size([3, 5])
```



Paper Discussion Time!

<http://cs.stanford.edu/~quocle/tutorial1.pdf>

Read through, and we will reconvene in ~5minutes

****Feel free to start discussing with your peers if you finish beforehand!**

Real World Applications

Image Identification (CNNs)
(RNNs)



Natural Language Processing



Additional Resources

Andrew Ng's deeplearning.ai course (highly recommend!):

<https://www.coursera.org/learn/neural-networks-deep-learning/lecture/Cuf2f/welcome>

Segway for Next Workshop: PyTorch 60minute Blitz (with CNNs!):

http://pytorch.org/tutorials/beginner/deep_learning_60min_blitz.html

Challenge Yourself: Efficient Backprop by Yann LeCun

<http://yann.lecun.com/exdb/publis/pdf/lecun-98b.pdf>