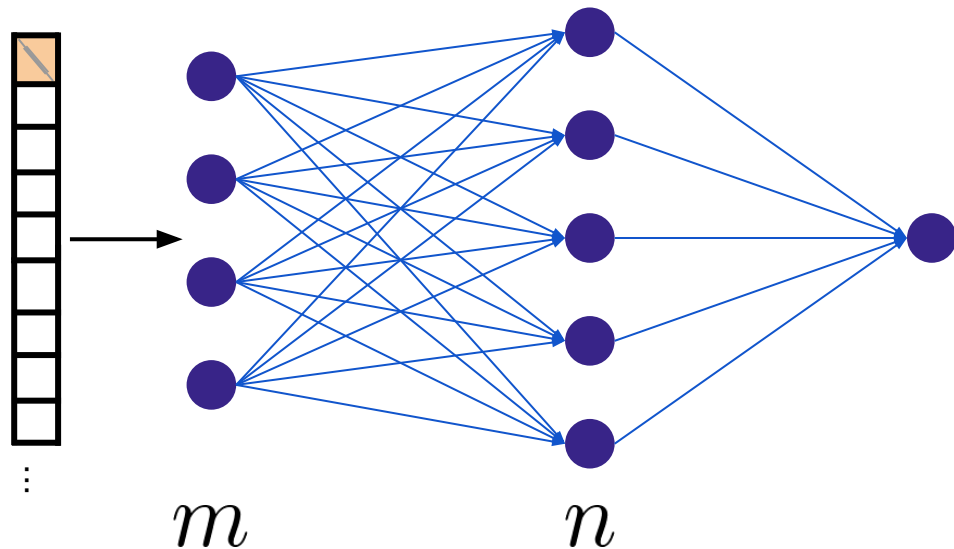


Transfer Learning

BOSTON UNIVERSITY
MACHINE INTELLIGENCE
COMMUNITY

Duy Nguyen, Justin Chen
Oct. 31, 2017

CNN: Fully-connected Layer



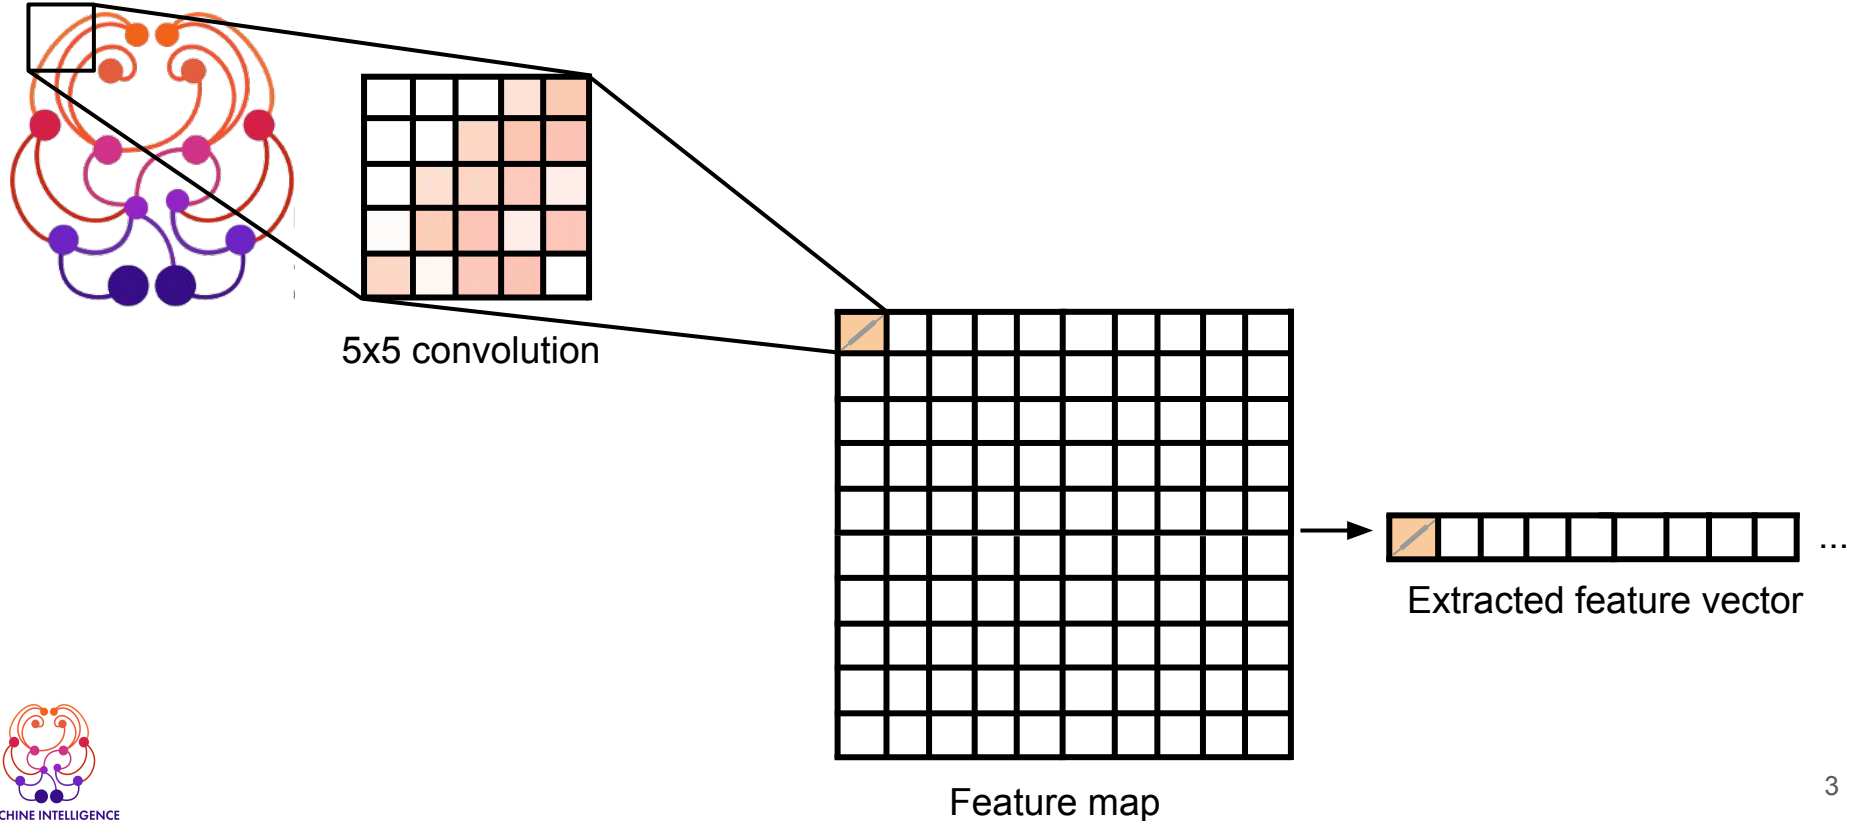
$$\begin{bmatrix} \theta_{11}^1 & \theta_{12}^1 & \theta_{13}^1 & \dots & \theta_{1m}^1 \\ \theta_{21}^1 & \theta_{22}^1 & \theta_{23}^1 & \dots & \theta_{2m}^1 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ \theta_{n1}^1 & \theta_{n2}^1 & \theta_{n3}^1 & \dots & \theta_{nm}^1 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_m \end{bmatrix} = \begin{bmatrix} a_1 \\ a_2 \\ \vdots \\ a_n \end{bmatrix}$$

$$\begin{bmatrix} \theta_{11}^2 & \theta_{12}^2 & \theta_{13}^2 & \dots & \theta_{1n}^2 \end{bmatrix} \sigma \left(\begin{bmatrix} a_1 \\ a_2 \\ \vdots \\ a_n \end{bmatrix} \right) = \begin{bmatrix} h_1 \end{bmatrix}$$

Ignoring bias values for simplification

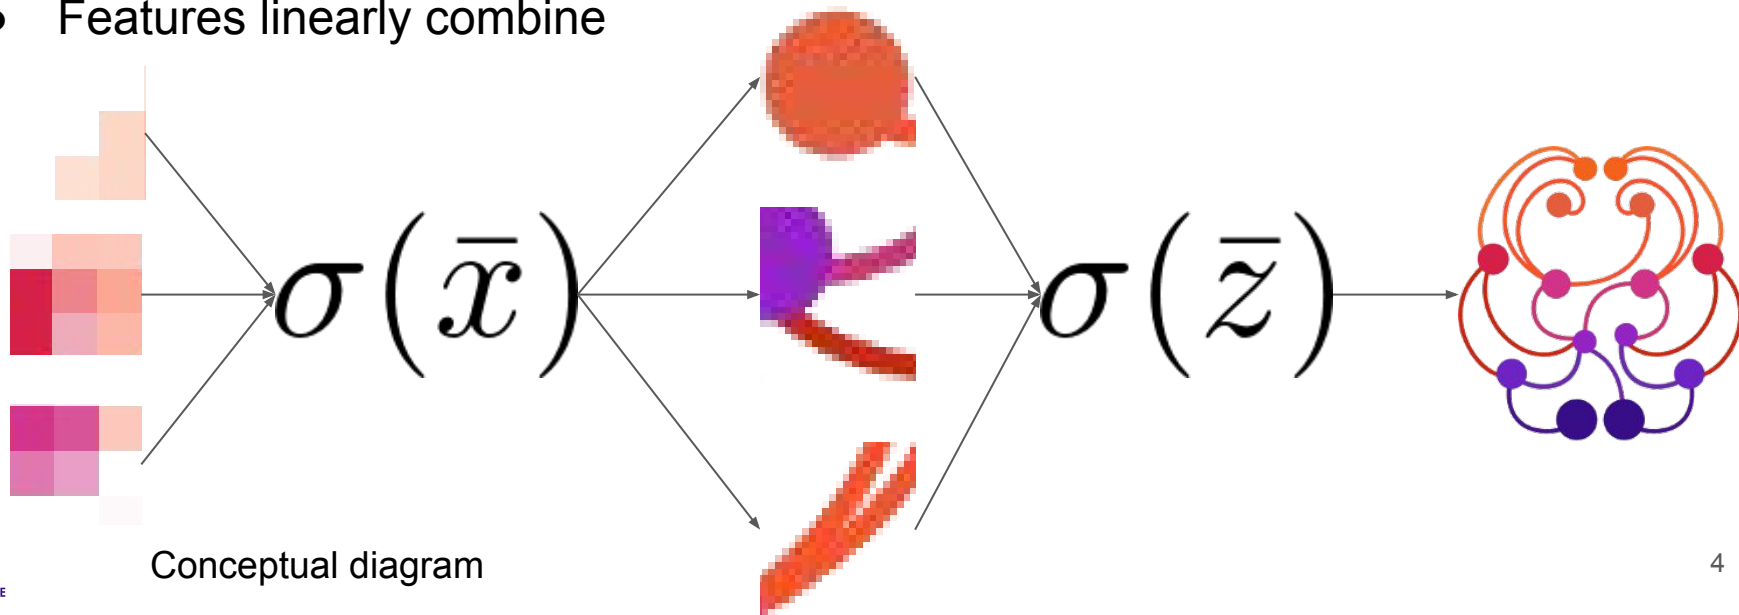
In this example, we're doing regression

CNN: Automatic Feature Extraction



Feature Hierarchy

- Learned features become progressively more complex throughout the network
- Earlier layers contain more primitive features
- Features linearly combine



Primitive features

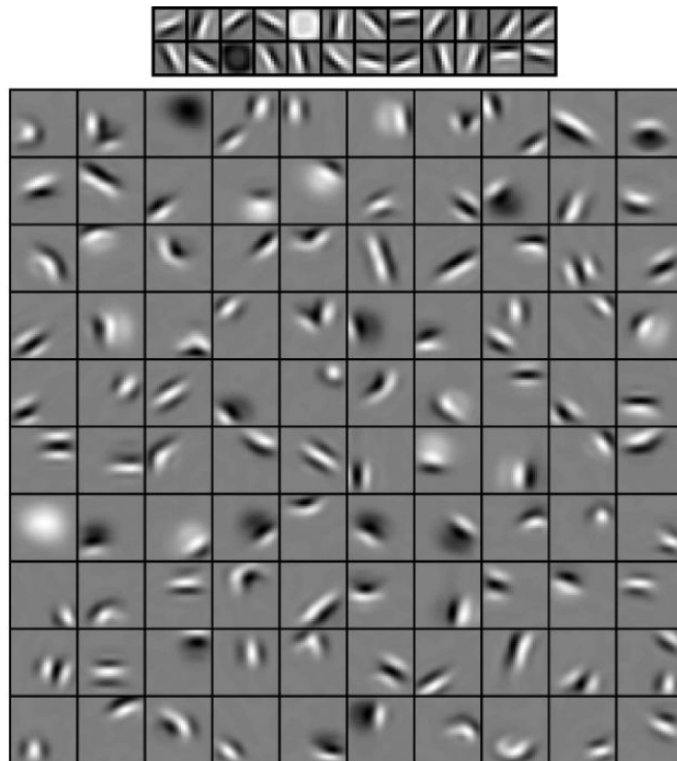
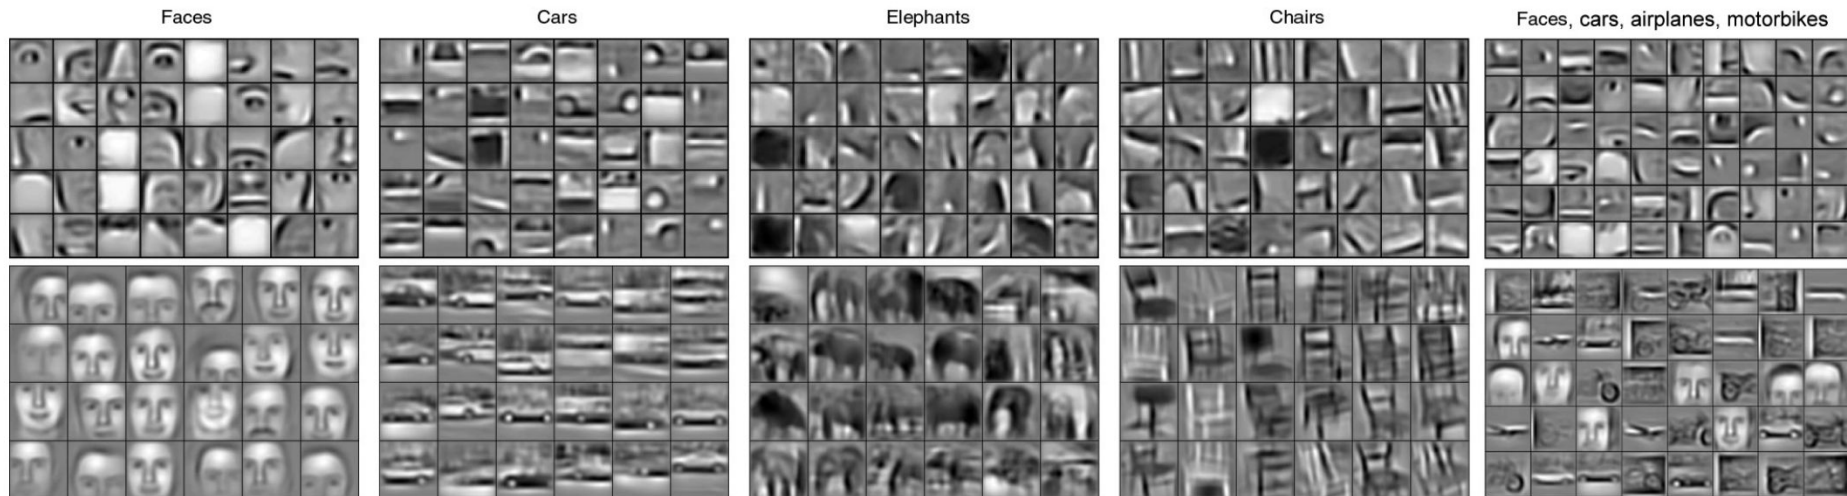


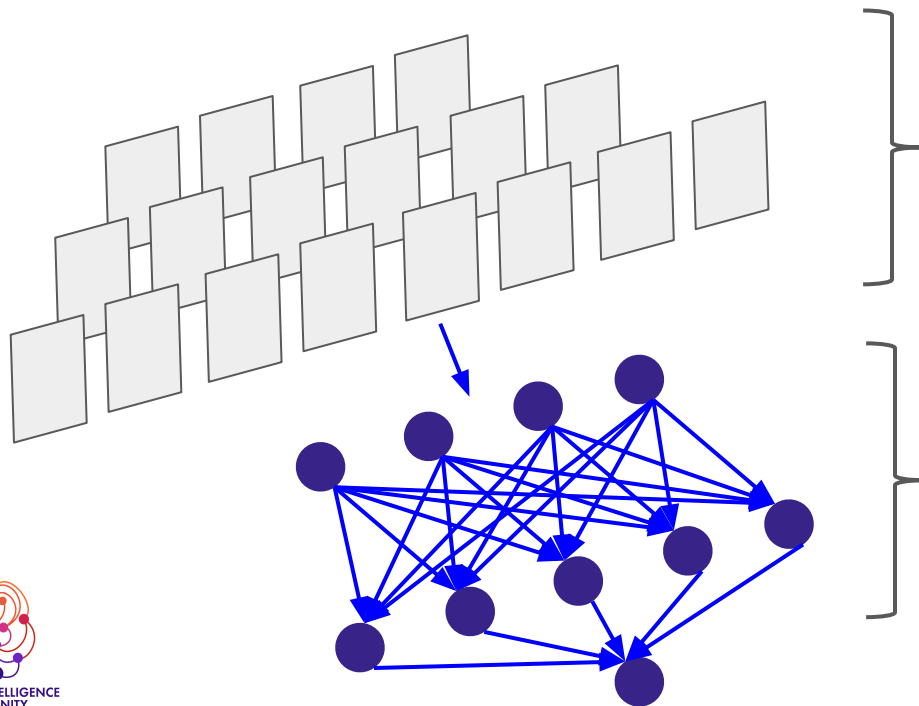
Image from Lee et al. 2011

Complex features



How to reuse learned models?

- Would be useful to be able to reuse learned parameters and learned features



- Trained feature extractor contains features common among domain
 - Reuse this part
-
- Trained fully-connected network contains features specific to dataset
 - Can potentially reuse weights in latter layers as well

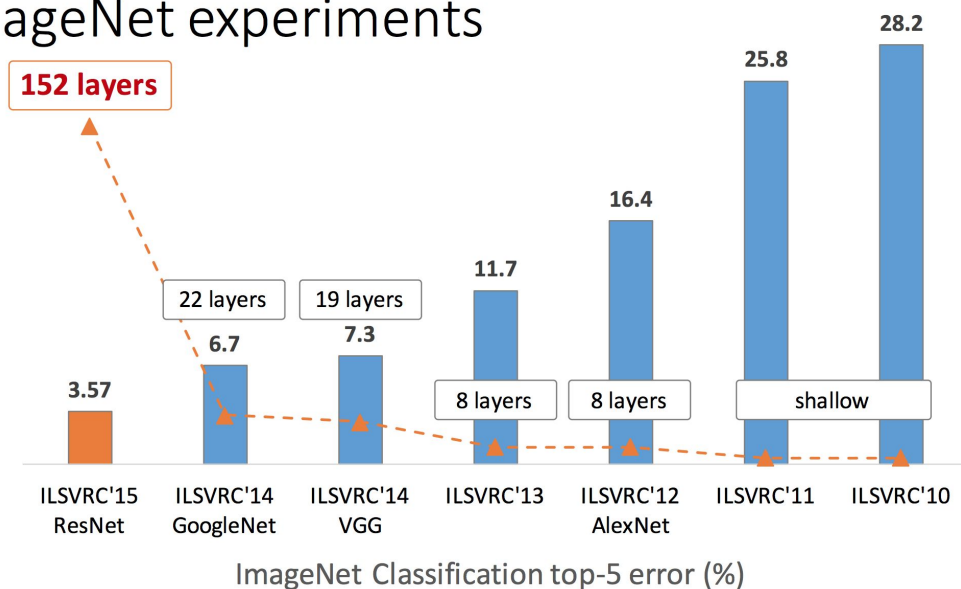
Deep learning so far ...

- Convolutional neural networks allows us to capture compositional data and convert it into a hierarchical representation.
 - This approach have prove quite effective in classification task, especially in computer vision.
- While CNN is effective in computer vision task, it is also requires a lot of resources.
- CNN, Neural Networks, Backpropagation are not new. All of these techniques have been around since the 1980s. But why it's popular now?

ImageNet

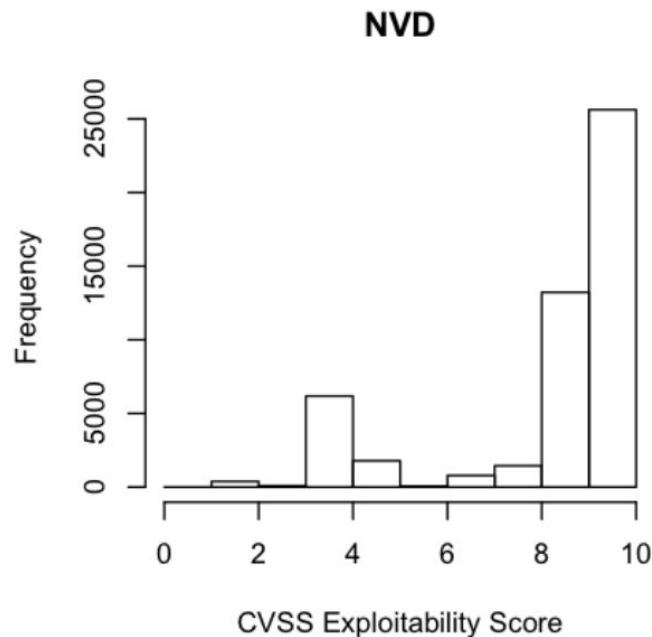
- Both a collection of 14+ millions images that are categorised **by human** into roughly 20 thousands categories. This dataset also fuels the ImageNet challenge.

ImageNet experiments



The domain of data

- Finding large amount of data for every task we want our machine learning model to perform is hard.
 - e.g. doing machine translation on rare language pair.
 - e.g. data about cyber vulnerabilities exploits



The domain of data

- Sometime our input data varies.
 - e.g. speech recognition is challenging because every person's speech is a little bit different. Furthermore, there are different accent of the same language, thus one model for one language would not work well.

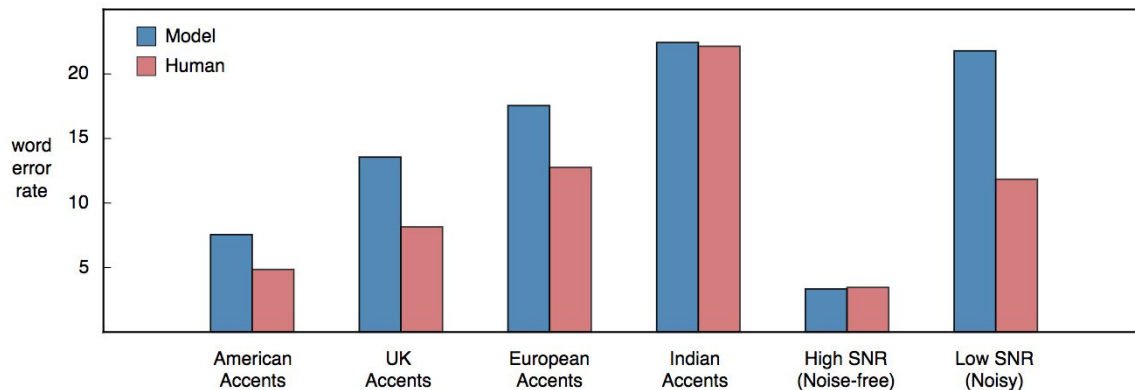
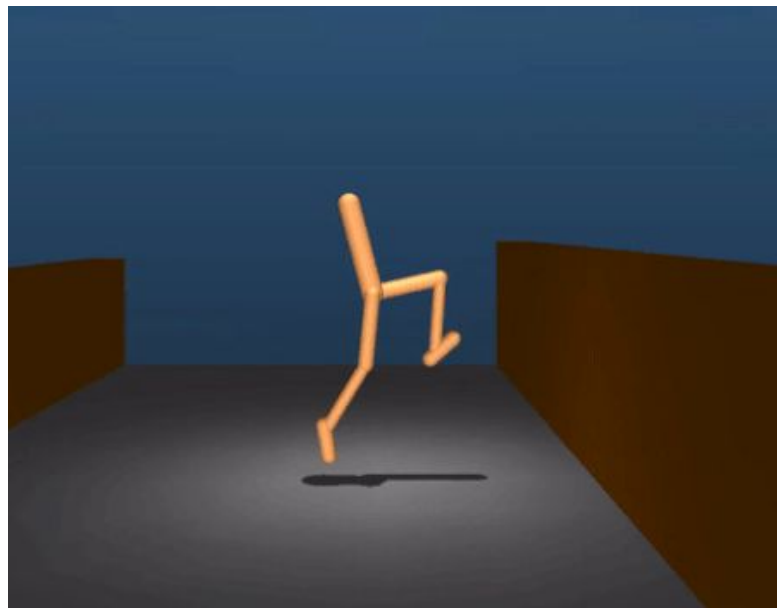


Chart from [Awni Hannun](#)

Learning from simulation

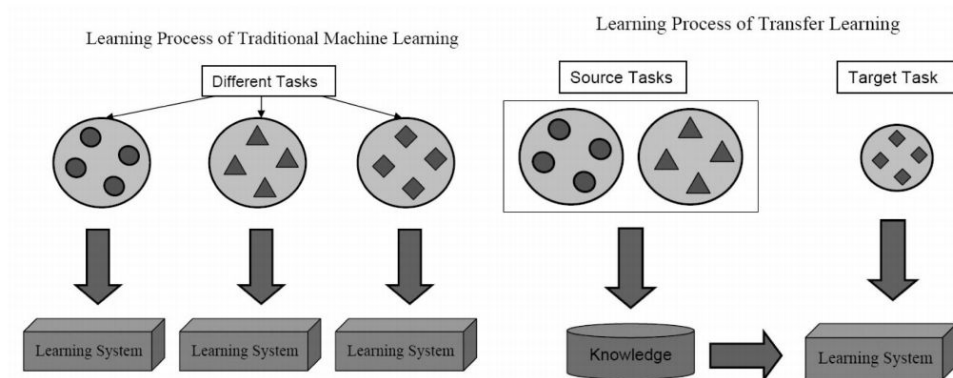
- Use a simulation of the real world to train a model.
 - Faster training time
 - Less costly
 - Able to repeat multiple time
- But the data from the simulation is still not the same data from the real world.
 - We need the ability to adapt the model from using the data in the simulated world to the data



DeepMind's simulated environment

So, transfer learning

- Utilise the feature extraction property of Convolutional Neural Networks
 - Reduce training time and computational cost.
 - Reduce the amount of data used.
- Improve performance with task where the data domain varies.
 - In general, instead of mapping data to output, we distill some “knowledge” about the data and use that to make a better machine learning model.



More formally

$$\mathcal{D} = \{\mathcal{X}, P(X)\}$$

Domain Feature space Marginal probability distribution

When $D_s \neq D_t$ or $T_s \neq T_t$,
transfer learning aims
to improve the target predictive function
(the conditional probability distribution)
in D_t using knowledge from D_s and T_s

Task Output/Label space

$$\mathcal{T} = \{\mathcal{Y}, P(Y|X)\}$$

Conditional probability distribution
(this is what's usually learnt)

Different scenarios of transfer learning

The domains are different

$$\mathcal{X}_S \neq \mathcal{X}_T$$

The feature spaces of the data are different from each other.
E.g. different languages.

$$P(X_S) \neq P(X_T)$$

The marginal probability distributions of the data are different.
E.g. different topic of document.

The tasks are different

$$\mathcal{Y}_S \neq \mathcal{Y}_T$$

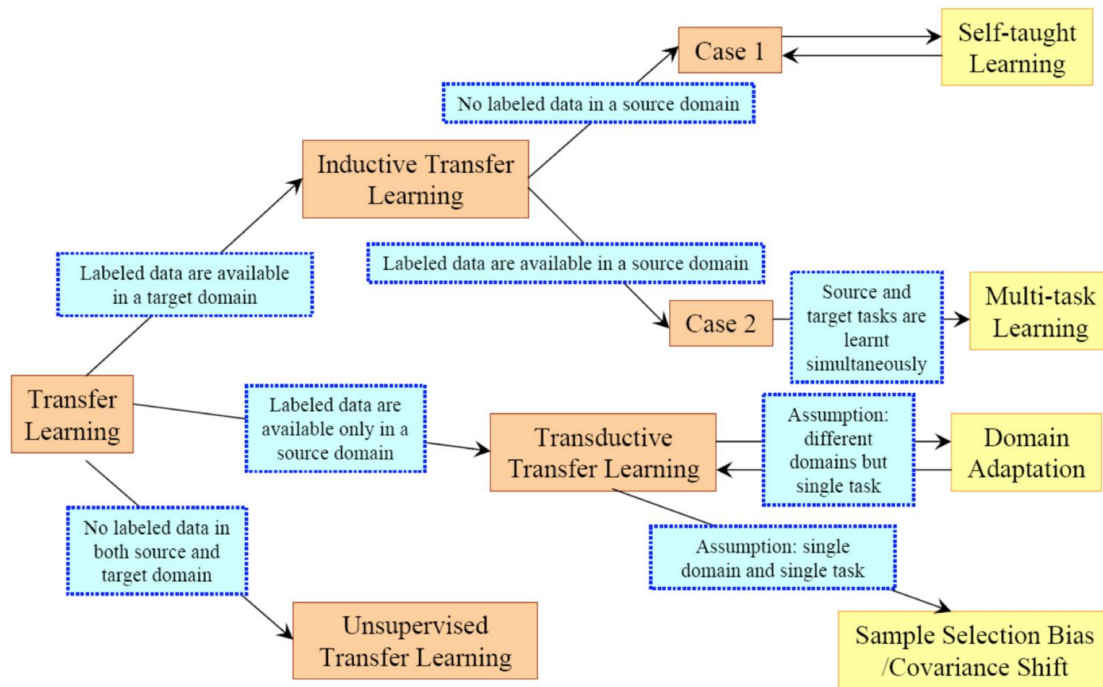
The label spaces are different.
E.g. classify the same document into different label depend on different task.

$$P(Y_S|X_S) \neq P(Y_T|X_T)$$

The learned functions are different.
E.g. the source and target distribution of the data is unbalanced compared to their class.

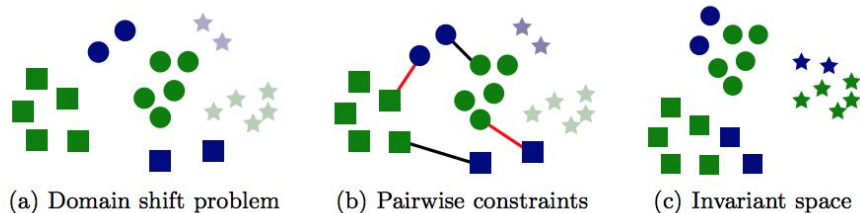
Different setting of transfer learning

- Both domain and task are different -> **Unsupervised Transfer Learning.**
- Domains are different -> **Transductive Transfer Learning**
- Tasks are different -> **Inductive Transfer Learning**



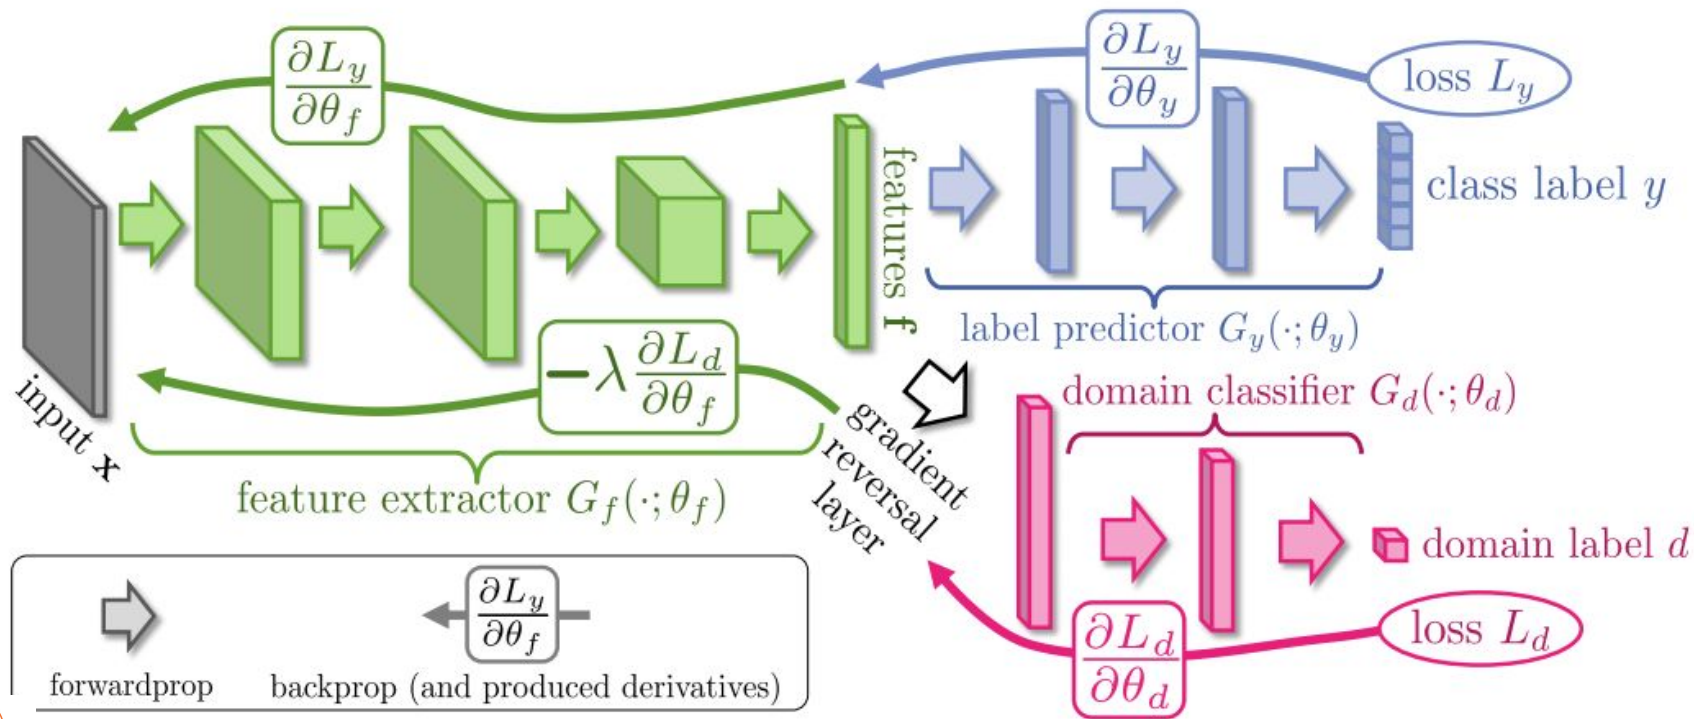
Domain Adaptation

- We are learning the same task but on different domain of data.
- To make our model adapt to the change in domain, we could make the two domain “closer” to each other.
 - Transform data from the domains into an invariant space ([K. Saenko et al. 2010](#))

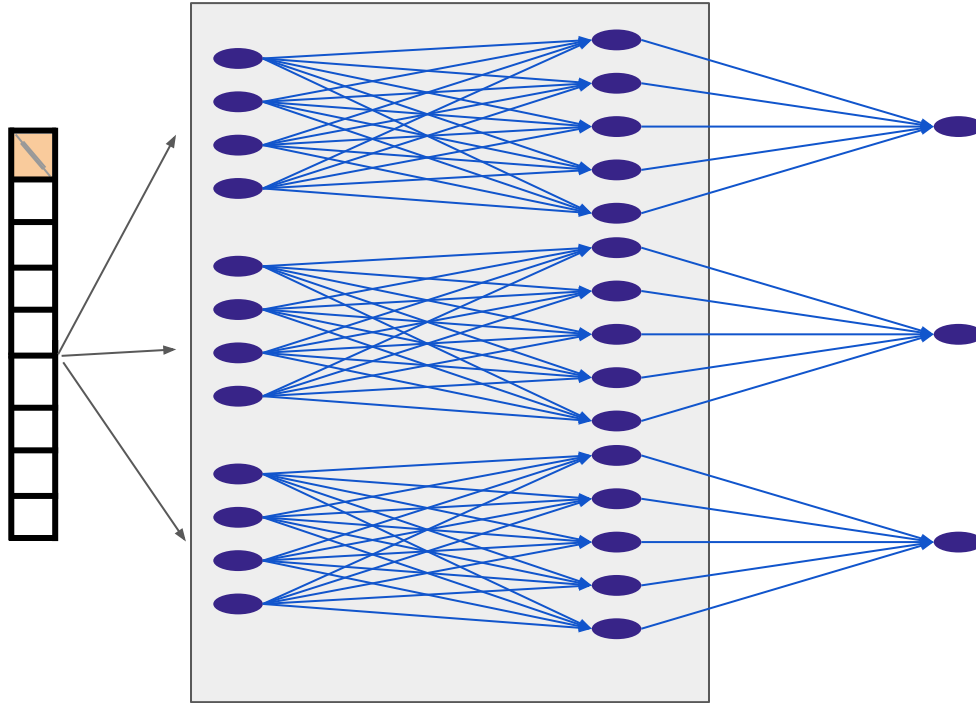


- Make it an objective of the model to “confuse” the domains of the data ([Ganin 2015](#))

Domain confusion using backpropagation



Multi-task learning



- Instead of training a model for each task, let's train one model for all the task and share the parameters for the models.
- This take advantage of the hierarchical representation of the data through the network.

Few-shot, One-shot, Zero-shot learning

- Data is expensive. Can we use transfer learning to learn with less data?
- **One-shot** learning means that the model have one example of the class before it has to make the prediction and **zero-shot** learning is making prediction on brand new data.

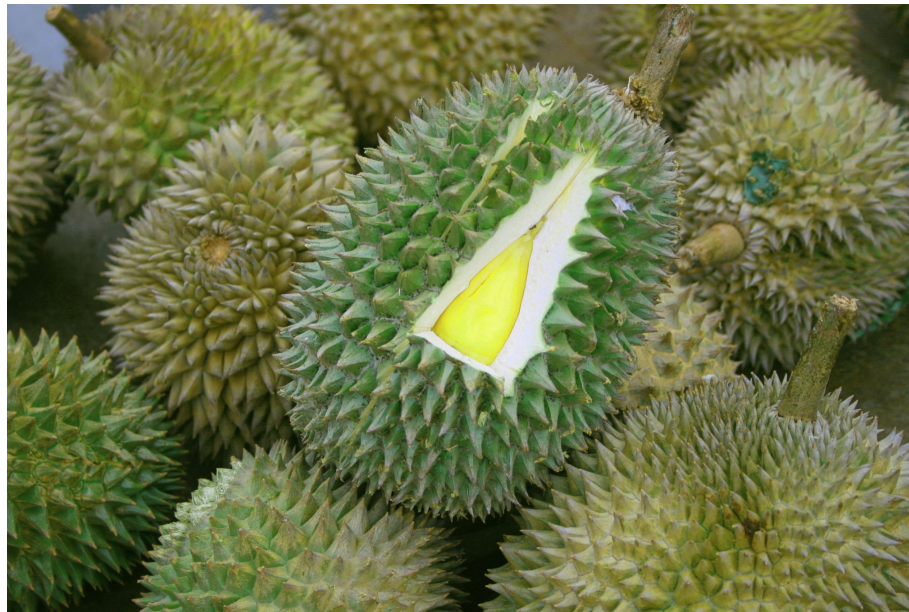
Example of few shot learning in human



- If you are not familiar with this spiky fruit on the left here, it's a durian.
- Human has the ability to learn/understand concepts, ideas, or objects from very small “training” example.

Example of few shot learning in human

- So if you see the image of this spiky fruit again, you would easily be able to recognise that's it's a durian.
 - In essence, you only require one or few training example to “learn” new thing.

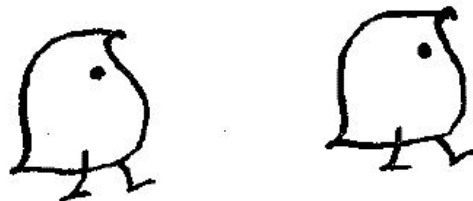


Example of zero shot learning in human

- Furthermore, human can generalise from other knowledge.
 - The wug test developed by Berko Gleason (1958) show that very young child can apply morphological rule to unknown words.
 - Here, we don't know the plural form of the word “wug”, but we are able to formulate it anyway.



THIS IS A WUG.



NOW THERE IS ANOTHER ONE.

THERE ARE TWO OF THEM.

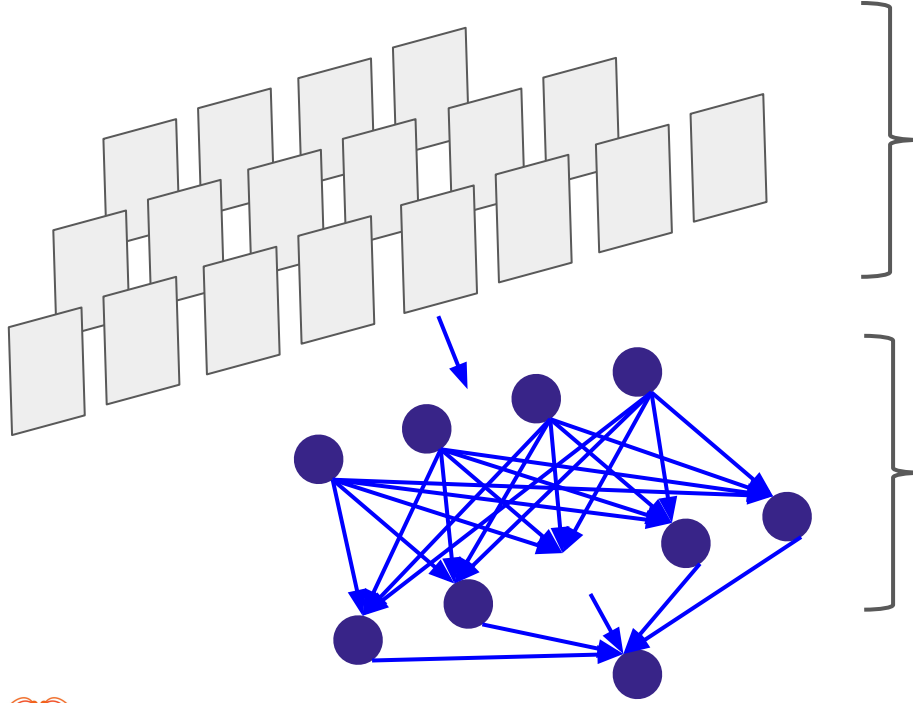
THERE ARE TWO _____.

Word-embedding

- Turns words to vector representation.
- Models like *word2vec* exploit the fact that semantically similar words tend to be around similar words.
- Vector representation have semantics value.



Pre-training and fine-tuning



- **Fine-tuning** is when you have a different but related dataset and you want to train the whole model to better fit what you have.
- **Pre-train** models have all the weights already trained. Just make a new Fully Connected layer at the end and retrain the layer.

Pretrained models with **PYTORCH**

```
1 import torchvision.models as models
2 resnet18 = models.resnet18(pretrained=True)
3 alexnet = models.alexnet(pretrained=True)
4 squeezenet = models.squeezenet1_0(pretrained=True)
5 vgg16 = models.vgg16(pretrained=True)
6 densenet = models.densenet161(pretrained=True)
7 inception = models.inception_v3(pretrained=True)
```

```
2 for param in model_conv.parameters():
3     param.requires_grad = False
```

```
1 num_fters = model_conv.fc.in_features
2 model_conv.fc = nn.Linear(num_fters, 3)
```

- Import and instantiate pretrained model
- Remove the output layer of model
- Append a randomly initialized classifier and fine-tune

[Try it out!](#)

References & Further Reading

1. http://machinelearning.wustl.edu/mlpapers/paper_files/ICML2011Glorot_342.pdf
2. <http://www.deeplearningbook.org/contents/representation.html>
3. <http://anthropology.uwo.ca/faculty/creider/027/wugs.pdf>
4. <https://www.cs.princeton.edu/~rajeshr/papers/cacm2011-researchHighlights-convDBN.pdf>
5. <http://runder.io/transfer-learning>
6. <http://ieeexplore.ieee.org/document/5288526/>

Shoutout to our Sponsors



Boston University Computer Science



Boston University SPARK!



Boston University Software
Application and Innovation Lab

Upcoming Events

MIC Paper signup: <https://goo.gl/iAm6TL>
BUMIC Projects signup: <https://goo.gl/GmP9oK>

BUMIC paper discussion:

Paper: **Decoupled Neural Interfaces using Synthetic Gradients**

Location: Fishbowl Conference Room

Date: 11.06.17 Time: 7 PM

Next workshop:

Topic: Sequential Data

Location: BU Hariri Seminar Room

Date: 11.07.17 Time: 7 PM